

Machine Learning Module 1

Introduction & Concept Learning

Educational Slide Deck

Slide 1: The \$2 Million Question

Why Can't Computers Learn Like We Do?

Think About This:

- Netflix recommends shows you'll love—how does it know?
- Gmail blocks spam emails automatically—who taught it?
- Self-driving cars recognize pedestrians—but never took a driving lesson

The Business Impact:

- Companies lose \$62 billion annually due to poor customer understanding
- Machine Learning reduces fraud detection costs by 40%
- Personalization engines increase revenue by 15-20%

The Challenge: Can we teach computers to improve from experience without explicitly programming every scenario?

Key Insight: The ability to learn from data is transforming industries—from healthcare to finance to entertainment.

Slide 2: Learning Like Humans Do

Solving Problems Through Experience

The Human Way - Problem: Solve $2x + 3 = 9$

Trial 1: "Let me try $x = 2$ "

- Result: $2(2) + 3 = 7$
- Error: Off by 2
- Learning: "x needs to be bigger"

Trial 2: "Now try $x = 3$ "

- Result: $2(3) + 3 = 9$ ✓

- Success!
- Memory: "When I see similar patterns, I know what to try"

Neural Network Parallel:

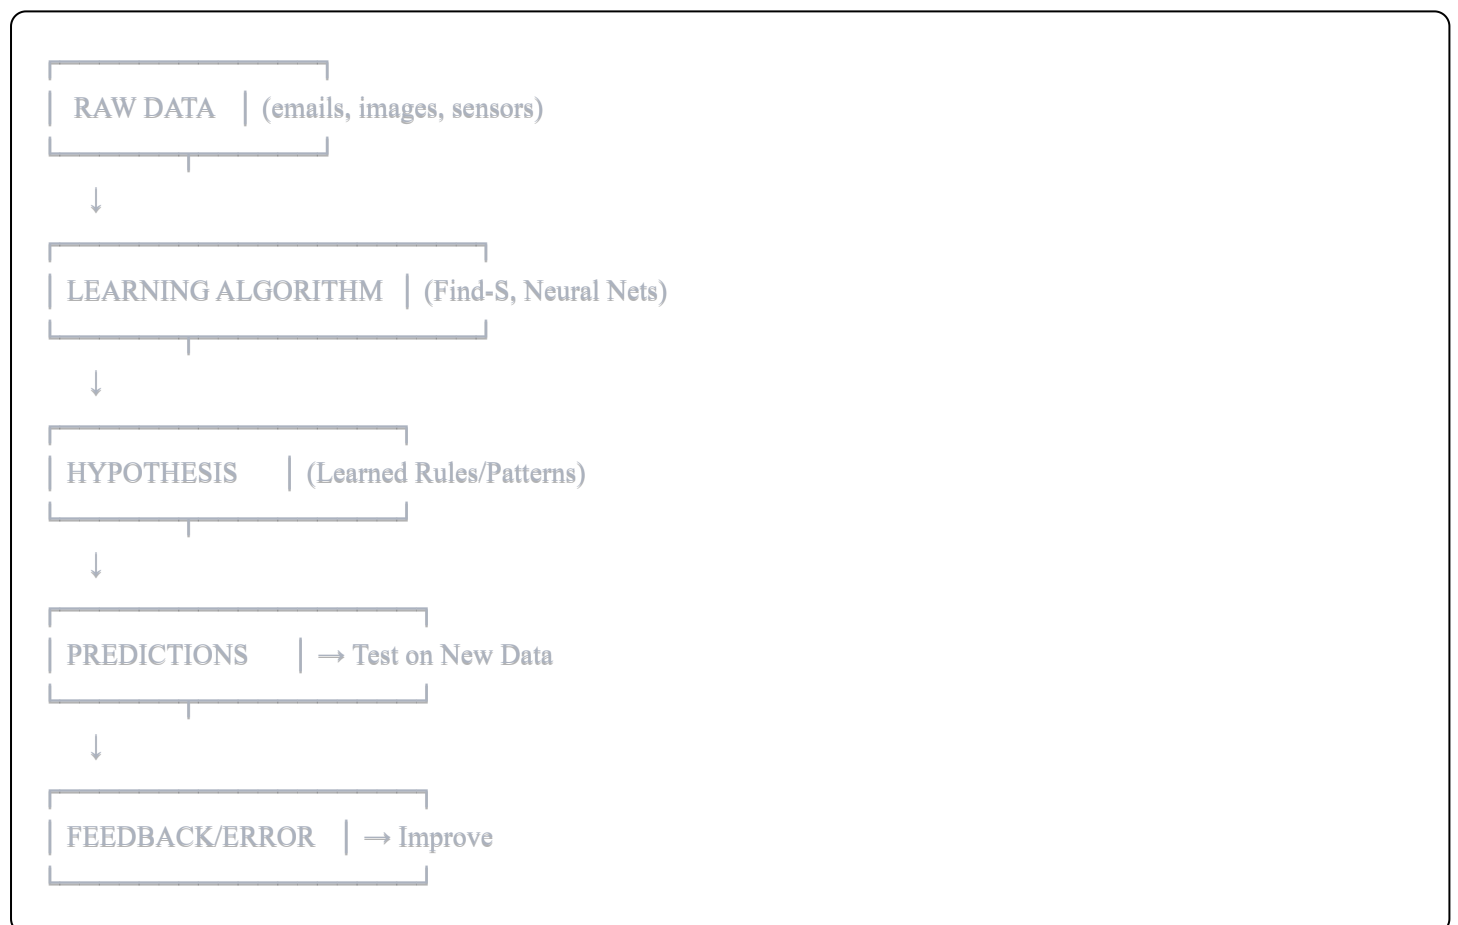
- Trial = Iteration/Epoch
- Error = Loss Function
- Adjustment = Gradient Descent
- Memory = Learned Weights
- Past Experience = Bias Term

Connection: ML algorithms use these same principles!

Slide 3: The ML Learning Cycle

How Machines Learn - Visual Overview

[DIAGRAM: Circular Flow]



Key Components:

- Data feeds the algorithm
- Algorithm searches for patterns

- Creates a hypothesis (model)
 - Tests and improves iteratively
-

Slide 4: Learning Objectives

What You'll Be Able To Do

By the end of this module, you will:

1. **Define** well-posed learning problems with task, performance, and experience
2. **Design** the four key components of any learning system
3. **Explain** concept learning as a search through hypothesis space
4. **Apply** the Find-S algorithm to find maximally specific hypotheses
5. **Implement** the Candidate-Elimination algorithm using version spaces
6. **Analyze** the role of inductive bias in learning

Career Impact: These fundamentals form the foundation for understanding modern ML frameworks like TensorFlow, PyTorch, and scikit-learn.

Slide 5: What is a Well-Posed Learning Problem?

The Three Essential Elements

Definition: A computer program learns from experience E with respect to some task T and performance measure P , if its performance at T (measured by P) improves with experience E .

The Three Components:

1. Task (T) - What are we trying to do?

- Classify emails as spam/not spam
- Predict house prices
- Play chess

2. Performance Measure (P) - How do we measure success?

- Accuracy percentage
- Mean squared error
- Win rate

3. Experience (E) - What data do we learn from?

- Labeled examples (supervised)
- Unlabeled data (unsupervised)
- Feedback from actions (reinforcement)

Think About: Banking fraud detection

- T: Classify transactions as fraudulent or legitimate
 - P: Percentage correctly classified
 - E: Database of past transactions with fraud labels
-

Slide 6: Real-World Example - Email Spam Filter

Applying the Well-Posed Framework

Task (T): Classify emails as spam or not spam

Performance Measure (P):

- Accuracy: % of emails correctly classified
- False Positive Rate: % of legitimate emails marked as spam (Critical!)
- False Negative Rate: % of spam emails that get through

Experience (E):

- Thousands of labeled emails (spam/not spam)
- User feedback (marking emails as spam)
- Features: sender, subject keywords, links, attachments

Business Impact:

- Gmail blocks 99.9% of spam and phishing attempts
- Processes 100+ billion spam attempts daily
- False positive rate < 0.05% (critical for user trust)

Why This Matters: A well-defined problem is half solved!

Slide 7: Designing a Learning System - The Checkers Game

Four Key Design Choices

The Classic Example: Teaching a computer to play checkers

Step-by-Step Design:

Choice 1: Determine the Training Experience

- Direct or Indirect feedback?
- Teacher control vs. Self-play?
- Representative distribution of examples?

Choice 2: Choose the Target Function

- What exactly should the system learn?
- ChooseMove: Board \rightarrow Move (too complex!)
- V: Board $\rightarrow \mathbb{R}$ (better - evaluation score)

Choice 3: Choose Representation

- How to represent the learned function?
- Linear combination? Decision tree? Neural network?

Choice 4: Choose Learning Algorithm

- How to fit the data to the representation?
- Minimize error? Gradient descent?

Key Insight: These four choices apply to ANY machine learning problem!

Slide 8: Checkers Example - Target Function Design

What Should the Computer Learn?

Option 1: ChooseMove Function

- Input: Current board state
- Output: Best legal move
- Problem: Too complex, difficult to learn directly!

Option 2: Board Evaluation Function V ✓

- Input: Current board state
- Output: Score (how good is this position?)
- $V(b) = +100$ if win, -100 if loss, 0 if draw
- For non-terminal: Estimate how likely to win from this state

Representation Choice - Linear Function:

$$V(b) = w_0 + w_1 \cdot x_1 + w_2 \cdot x_2 + w_3 \cdot x_3 + w_4 \cdot x_4 + w_5 \cdot x_5 + w_6 \cdot x_6$$

Board Features:

- x_1 = # black pieces on board
- x_2 = # red pieces on board
- x_3 = # black kings
- x_4 = # red kings
- x_5 = # black pieces threatened
- x_6 = # red pieces threatened

The Learning Task: Find weights $w_0 \dots w_6$ that best predict game outcomes!

Slide 9: Quick Check - Understanding the Basics

Test Your Knowledge

Question 1 (Conceptual): Which component determines "how we measure success" in a learning problem?

- a) Task
- b) Performance Measure ✓
- c) Experience
- d) Hypothesis

Question 2 (Predictive): If we increase training data from 1,000 to 10,000 examples, what happens?

- a) Task changes
- b) Performance typically improves ✓
- c) Algorithm changes
- d) Nothing changes

Question 3 (Practical): For a house price prediction system, which is the best performance measure?

- a) Number of houses in dataset
- b) Mean Absolute Error (\$ difference) ✓
- c) Number of features
- d) Training time

Discuss: Why is defining the right performance measure critical for business success?

Slide 10: Concept Learning - The Core Problem

Learning Boolean-Valued Functions

What is Concept Learning? Inferring a boolean-valued function from training examples of its input and output.

Everyday Example: Learning "days Aldo enjoys water sports"

Attributes describe each day:

- Sky: Sunny, Cloudy, Rainy
- AirTemp: Warm, Cold
- Humidity: Normal, High
- Wind: Strong, Weak
- Water: Warm, Cool
- Forecast: Same, Change

Target Concept: EnjoySport: Yes or No

Training Data:

Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoySport
Sunny	Warm	Normal	Strong	Warm	Same	Yes
Sunny	Warm	High	Strong	Warm	Same	Yes
Rainy	Cold	High	Strong	Warm	Change	No
Sunny	Warm	High	Strong	Cool	Change	Yes

The Challenge: Can you predict if Aldo will enjoy water sports on a new day?

Slide 11: Hypothesis Representation

Describing Concepts with Constraints

How do we represent hypotheses?

Each hypothesis = conjunction of constraints on attributes

Three Types of Constraints:

1. "?" (**Any value**) - This attribute can be anything
2. **Specific value** - Must match exactly (e.g., "Warm")
3. "∅" (**No value**) - Impossible to satisfy (all negative)

Example Hypotheses:

$h_1 = \langle \text{Sunny}, ?, ?, \text{Strong}, ?, ? \rangle$

- Means: Sky must be Sunny, Wind must be Strong, others don't matter

$h_2 = \langle \text{Sunny}, ?, ?, ?, ?, ? \rangle$

- Means: Only Sky = Sunny matters

$h_3 = \langle ?, \text{Warm}, \text{Normal}, \text{Strong}, \text{Warm}, \text{Same} \rangle$

- Means: All six attributes must match exactly

Most General Hypothesis: $\langle ?, ?, ?, ?, ?, ? \rangle$ (all instances positive)

Most Specific Hypothesis: $\langle \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset \rangle$ (all instances negative)

Hypothesis Space Size: $3 \times 2 \times 2 \times 2 \times 2 \times 2 = 96$ syntactically distinct

- Plus 1 empty hypothesis = 97 total
 - But only 973 semantically distinct concepts possible!
-

Slide 12: The General-to-Specific Ordering

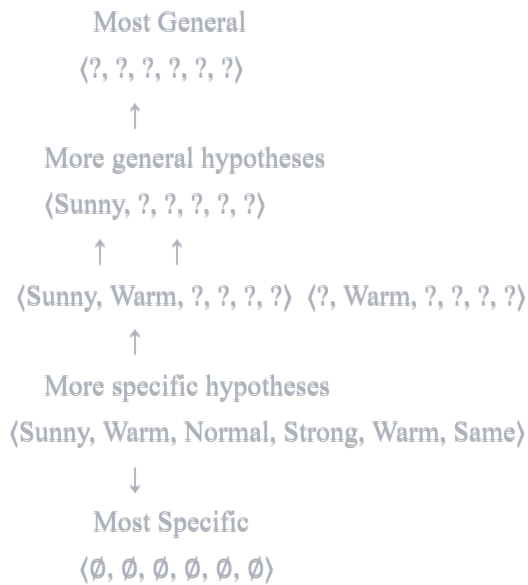
Structure of the Hypothesis Space

Key Idea: Hypotheses can be ordered by specificity

Definition: Hypothesis h_j is **more general than or equal to** h_k (written $h_j \geq_m h_k$) if:

- Every instance classified positive by h_k is also classified positive by h_j

Visual Representation:



Why This Matters:

- Provides structure to search through hypothesis space
- Enables efficient learning algorithms
- Basis for Find-S and Candidate-Elimination

Slide 13: Find-S Algorithm - Finding Maximally Specific Hypothesis

Step-by-Step Learning Process

Algorithm Goal: Find the most specific hypothesis that fits all positive examples

The Find-S Algorithm:

1. Initialize h to the most specific hypothesis $\langle \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset \rangle$
2. For each positive training example x :
 - For each attribute constraint a_i in h :
 - If constraint a_i is satisfied by x :
 - Do nothing
 - Else:
 - Replace a_i in h by the next more general constraint that is satisfied by x
3. Output hypothesis h

Key Properties:

- Only considers positive examples

- Moves from specific to general
- Guaranteed to find maximally specific hypothesis
- Ignores negative examples (limitation!)

Python Code Example:

```
python

def find_s(examples):
    # Initialize to most specific
    h = ['?', '?', '?', '?', '?', '?']

    for x, label in examples:
        if label == 'Yes': # Only positive examples
            for i in range(len(h)):
                if h[i] == '?':
                    h[i] = x[i] # First positive example
            elif h[i] != x[i]:
                h[i] = '?' # Generalize if different

    return h
```

Slide 14: Find-S Example Trace

Learning from the EnjoySport Data

Training Examples:

1. ⟨Sunny, Warm, Normal, Strong, Warm, Same⟩ → **Yes**
2. ⟨Sunny, Warm, High, Strong, Warm, Same⟩ → **Yes**
3. ⟨Rainy, Cold, High, Strong, Warm, Change⟩ → **No**
4. ⟨Sunny, Warm, High, Strong, Cool, Change⟩ → **Yes**

Step-by-Step Execution:

Initial: $h_0 = \langle ?, ?, ?, ?, ?, ? \rangle$

After Example 1 (+): $h_1 = \langle \text{Sunny, Warm, Normal, Strong, Warm, Same} \rangle$

- First positive example, copy all attributes

After Example 2 (+): $h_2 = \langle \text{Sunny, Warm, ?, Strong, Warm, Same} \rangle$

- Humidity differs (Normal vs High) → generalize to "?"

After Example 3 (-): $h_3 = \langle \text{Sunny, Warm, ?, Strong, Warm, Same} \rangle$

- Negative example, no change (Find-S ignores negatives)

After Example 4 (+): $h_4 = \langle \text{Sunny, Warm, ?, Strong, ?, ?} \rangle$

- Water differs (Warm vs Cool) \rightarrow "?"
- Forecast differs (Same vs Change) \rightarrow "?"

Final Hypothesis: $\langle \text{Sunny, Warm, ?, Strong, ?, ?} \rangle$

Interpretation: Aldo enjoys water sports on Sunny, Warm days with Strong wind!

Slide 15: Limitations of Find-S & Version Spaces

Why We Need Something Better

Find-S Problems:

1. Ignores Negative Examples

- Can't detect when hypothesis is too general
- Might classify negative examples as positive

2. No Way to Know if Converged

- Did we find the correct concept?
- Are there other equally good hypotheses?

3. Can't Handle Inconsistent Data

- What if training examples have noise/errors?
- No mechanism to detect contradictions

Solution: Version Spaces 

Key Idea: Instead of a single hypothesis, maintain the set of ALL hypotheses consistent with training data!

Version Space Definition: The subset of all hypotheses from H that are consistent with the observed training examples D .

Mathematical Notation: $VS_{h,D} = \{h \in H \mid \text{Consistent}(h, D)\}$

Representation: Rather than list all consistent hypotheses (could be thousands!), represent by:

- **S (Specific boundary):** Most specific consistent hypotheses
 - **G (General boundary):** Most general consistent hypotheses
-

Slide 16: Candidate-Elimination Algorithm

Maintaining Version Spaces Efficiently

The Algorithm:

Initialize:

$G \leftarrow \{\langle ?, ?, ?, ?, ?, ? \rangle\}$ (most general hypothesis)

$S \leftarrow \{\langle \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset \rangle\}$ (most specific hypothesis)

For each training example $d = \langle x, c(x) \rangle$:

If d is a POSITIVE example:

- Remove from G any hypothesis inconsistent with d
- For each s in S not consistent with d :
 - Remove s from S
 - Add all minimal generalizations h of s :
 - where h is consistent with d
 - and some member of G is more general than h
- Remove from S any hypothesis more general than another in S

If d is a NEGATIVE example:

- Remove from S any hypothesis inconsistent with d
- For each g in G not consistent with d :
 - Remove g from G
 - Add all minimal specializations h of g :
 - where h is consistent with d
 - and some member of S is more specific than h
- Remove from G any hypothesis less general than another in G

Key Insight: Version space is fully characterized by S and G boundaries!

Slide 17: Candidate-Elimination Example

Learning EnjoySport Step-by-Step

Training Data:

1. $\langle \text{Sunny, Warm, Normal, Strong, Warm, Same} \rangle \rightarrow \text{Yes}$
2. $\langle \text{Sunny, Warm, High, Strong, Warm, Same} \rangle \rightarrow \text{Yes}$
3. $\langle \text{Rainy, Cold, High, Strong, Warm, Change} \rangle \rightarrow \text{No}$
4. $\langle \text{Sunny, Warm, High, Strong, Cool, Change} \rangle \rightarrow \text{Yes}$

Initial State:

- $S_0 = \{\langle \emptyset, \emptyset, \emptyset, \emptyset, \emptyset, \emptyset \rangle\}$
- $G_0 = \{\langle ?, ?, ?, ?, ?, ? \rangle\}$

After Example 1 (+):

- $S_1 = \{\langle \text{Sunny, Warm, Normal, Strong, Warm, Same} \rangle\}$
- $G_1 = \{\langle ?, ?, ?, ?, ?, ? \rangle\}$ (no change)

After Example 2 (+):

- $S_2 = \{\langle \text{Sunny, Warm, ?, Strong, Warm, Same} \rangle\}$ (generalized Humidity)
- $G_2 = \{\langle ?, ?, ?, ?, ?, ? \rangle\}$

After Example 3 (-): $\langle \text{Rainy, Cold, High, Strong, Warm, Change} \rangle$

- $S_3 = \{\langle \text{Sunny, Warm, ?, Strong, Warm, Same} \rangle\}$ (still consistent)
- $G_3 = \{\langle \text{Sunny, ?, ?, ?, ?, ?} \rangle, \langle ?, \text{Warm, ?, ?, ?, ?} \rangle, \langle ?, ?, ?, ?, ?, \text{Same} \rangle\}$ (G specialized to exclude negative example)

After Example 4 (+):

- $S_4 = \{\langle \text{Sunny, Warm, ?, Strong, ?, ?} \rangle\}$
- $G_4 = \{\langle \text{Sunny, ?, ?, ?, ?, ?} \rangle, \langle ?, \text{Warm, ?, ?, ?, ?} \rangle\}$

Final Version Space: All hypotheses between S_4 and G_4

Slide 18: Interview Question - Version Spaces

Testing Your Understanding

Junior Level: *"What is the difference between Find-S and Candidate-Elimination algorithms?"*

Expected Answer:

- Find-S maintains only one hypothesis (most specific)
- Candidate-Elimination maintains all consistent hypotheses via S and G boundaries
- Find-S ignores negative examples; C-E uses both positive and negative
- C-E can detect inconsistencies; Find-S cannot

Mid Level: *"Given this version space, classify the new instance $\langle \text{Sunny, Warm, Normal, Light, Warm, Same} \rangle$ "*

With $S = \{\langle \text{Sunny, Warm, ?, Strong, ?, ?} \rangle\}$ and $G = \{\langle \text{Sunny, ?, ?, ?, ?, ?} \rangle, \langle ?, \text{Warm, ?, ?, ?, ?} \rangle\}$

Expected Answer:

- Check against S: Wind is Light (not Strong) → **Not satisfied by S**
- Check against G: Both satisfied
- Result: **Cannot classify with confidence** (some hypotheses say yes, some say no)
- Need more training data to narrow version space

Senior Level: *"What happens to the version space size as we add more training examples? Can it increase?"*

Expected Answer:

- Version space can only shrink or stay same (monotonic)
 - Each example eliminates hypotheses
 - Cannot increase because we never add back eliminated hypotheses
 - Converges when $S = G$ (single hypothesis remains)
-

Slide 19: Inductive Bias - The Necessity of Assumptions

Why Learning Requires Bias

The Fundamental Problem:

Question: Can an unbiased learner generalize beyond training data?

Answer: NO! An unbiased learner cannot make inductive leaps.

Example - Unbiased Learner (Power Set Hypothesis Space):

Given 3 positive examples:

- ⟨Sunny, Warm, Normal, Strong, Warm, Same⟩ → Yes
- ⟨Sunny, Warm, High, Strong, Warm, Same⟩ → Yes
- ⟨Rainy, Cold, High, Strong, Cool, Change⟩ → No

With power set H: Every unobserved instance will be classified positive by exactly HALF the consistent hypotheses!

- Cannot make predictions with confidence
- Voting provides no information

The Futility of Bias-Free Learning: Without assumptions, every classification is equally likely!

Inductive Bias Definition: The set of assumptions a learner uses to predict outputs for inputs it has not encountered.

Candidate-Elimination Inductive Bias: *"The target concept can be represented as a conjunction of attribute constraints."*

Key Insight:

- Stronger bias → more assumptions → better generalization (if assumptions correct)
- Weaker bias → fewer assumptions → less generalization ability
- No bias → no generalization!

Trade-off: Bias allows learning but risks being wrong if assumptions don't hold.

Slide 20: Key Takeaways & Next Steps

Summary & Hands-On Assignment

Key Takeaways:

1. **Well-posed learning requires:** Task (T), Performance (P), Experience (E)
2. **Four design choices:** Training experience, target function, representation, learning algorithm
3. **Concept learning** is search through hypothesis space with general-to-specific ordering
4. **Find-S:** Finds maximally specific hypothesis (ignores negatives)
5. **Version Spaces:** All hypotheses consistent with data, represented by S and G boundaries
6. **Candidate-Elimination:** Efficiently maintains version space using both positive and negative examples
7. **Inductive bias is necessary:** Bias-free learning cannot generalize beyond training data

Hands-On Assignment (Due: Next Week):

Objective: Implement and test the Candidate-Elimination algorithm

Required Tasks:

1. Implement the Candidate-Elimination algorithm in Python
2. Apply it to the EnjoySport dataset (provided)
3. Visualize how S and G boundaries evolve with each example
4. Test classification on 5 new instances

Optional Challenge: Modify the algorithm to handle noisy data (contradictory examples)

Deliverables:

- Python code (well-commented)
- Report showing S/G evolution

- Analysis of 3 classification decisions

Next Module Preview:

- Decision Tree Learning
- Overfitting and generalization
- Practical ML with scikit-learn

Resources: Textbook Chapter 1-2, online notebooks at course website

End of Module 1 Slide Deck